

```
In [1]: import warnings
warnings.filterwarnings('ignore')

In [2]: from tensorflow.keras.layers import Dense, Flatten, Input, Lambda
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img
from glob import glob
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf

In [3]: import tensorflow as tf
import tensorflow_addons as tfa
from tensorflow import keras
from tensorflow.keras import layers

In [10]: input_shape = (256, 256, 3)
num_classes = 4

train_path = 'Brain_Tumor_MRI_Image_Dataset/Training'
test_path = 'Brain_Tumor_MRI_Image_Dataset/Testing'

In [11]: # Scaling all the images between 0 to 1

train_datagen = ImageDataGenerator(rescale = 1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=False)

# Performing only scaling on the test dataset

test_datagen = ImageDataGenerator(rescale=1./255)

In [12]: train_set = train_datagen.flow_from_directory(train_path,
                                                 target_size=(256,256),
                                                 batch_size=2,
                                                 class_mode = 'categorical')

test_set = test_datagen.flow_from_directory(test_path,
                                             target_size=(256,256),
                                             batch_size=2,
                                             class_mode='categorical')

Found 5712 images belonging to 4 classes.
Found 1311 images belonging to 4 classes.

In [3]: from keras import backend as K

def recall_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    recall = true_positives / (possible_positives + K.epsilon())
    return recall

def precision_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def f1_m(y_true, y_pred):
    precision = precision_m(y_true, y_pred)
    recall = recall_m(y_true, y_pred)
    return 2*((precision*recall)/(precision+recall+K.epsilon()))

In [4]: from tensorflow.keras.applications import Xception, NASNetLarge, NASNetMobile, VGG16, VGG19, InceptionV3, ResNet50, InceptionResNetV2

In [22]: from keras.metrics import Precision, Recall
from keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.models import Sequential, Model

learning_rate_reduction = ReduceLROnPlateau(
    monitor="val_accuracy", patience=3, verbose=1, factor=0.3, min_lr=0.0000001
)
early_stop = EarlyStopping(
    patience=10,
    verbose=1,
    monitor="val_accuracy",
    mode="max",
    min_delta=0.001,
    restore_best_weights=True,
)

In [65]: ML_Model = []
accuracy = []
precision = []
recall = []
f1score = []

#function to call for storing the results
def storeResults(model, a,b,c,d):
    ML_Model.append(model)
    accuracy.append(round(a, 3))
    precision.append(round(b, 3))
    recall.append(round(c, 3))
    f1score.append(round(d, 3))
```

## ViT Models

### SWIN

```
In [16]: patch_size = (2, 2) # 2-by-2 sized patches
dropout_rate = 0.03 # Dropout rate
num_heads = 8 # Attention heads
embed_dim = 64 # Embedding dimension
num_mlp = 256 # MLP Layer size
qkv_bias = True # Convert embedded patches to query, key, and values with a Learnable additive value
window_size = 2 # Size of attention window
shift_size = 1 # Size of shifting window
image_dimension = 256 # Initial image size

num_patch_x = input_shape[0] // patch_size[0]
num_patch_y = input_shape[1] // patch_size[1]

learning_rate = 1e-3
batch_size = 2
num_epochs = 40
validation_split = 0.1
weight_decay = 0.0001
label_smoothing = 0.1
```

```
In [17]: def window_partition(x, window_size):
    _, height, width, channels = x.shape
    patch_num_y = height // window_size
    patch_num_x = width // window_size
    x = tf.reshape(
        x, shape=(-1, patch_num_y, window_size, patch_num_x, window_size, channels)
    )
    x = tf.transpose(x, (0, 1, 3, 2, 4, 5))
    windows = tf.reshape(x, shape=(-1, window_size, window_size, channels))
    return windows

def window_reverse(windows, window_size, height, width, channels):
    patch_num_y = height // window_size
    patch_num_x = width // window_size
    x = tf.reshape(
        windows,
        shape=(-1, patch_num_y, patch_num_x, window_size, window_size, channels),
    )
    x = tf.transpose(x, perm=(0, 1, 3, 2, 4, 5))
    x = tf.reshape(x, shape=(-1, height, width, channels))
    return x

class DropPath(layers.Layer):
    def __init__(self, drop_prob=None, **kwargs):
        super(DropPath, self).__init__(**kwargs)
        self.drop_prob = drop_prob

    def call(self, x):
        input_shape = tf.shape(x)
        batch_size = input_shape[0]
        rank = x.shape.rank
        shape = (batch_size,) + (1,) * (rank - 1)
        random_tensor = (1 - self.drop_prob) + tf.random.uniform(shape, dtype=x.dtype)
        path_mask = tf.floor(random_tensor)
        output = tf.math.divide(x, 1 - self.drop_prob) * path_mask
        return output
```

```
In [18]: class WindowAttention(layers.Layer):
    def __init__(self, dim, window_size, num_heads, qkv_bias=True, dropout_rate=0.0, **kwargs):
        super(WindowAttention, self).__init__(**kwargs)
        self.dim = dim
        self.window_size = window_size
        self.num_heads = num_heads
        self.scale = (dim // num_heads) ** -0.5
        self.qkv = layers.Dense(dim * 3, use_bias=qkv_bias)
        self.dropout = layers.Dropout(dropout_rate)
        self.proj = layers.Dense(dim)

    def build(self, input_shape):
        num_window_elements = (2 * self.window_size[0] - 1) * (
            2 * self.window_size[1] - 1
        )
        self.relative_position_bias_table = self.add_weight(
            shape=(num_window_elements, self.num_heads),
            initializer=tf.initializers.Zeros(),
            trainable=True,
        )
        coords_h = np.arange(self.window_size[0])
        coords_w = np.arange(self.window_size[1])
        coords_matrix = np.meshgrid(coords_h, coords_w, indexing="ij")
        coords = np.stack(coords_matrix)
        coords_flatten = coords.reshape(2, -1)
        relative_coords = coords_flatten[:, :, None] - coords_flatten[:, None, :]
        relative_coords = relative_coords.transpose([1, 2, 0])
        relative_coords[:, :, 0] += self.window_size[0] - 1
        relative_coords[:, :, 1] += self.window_size[1] - 1
        relative_coords[:, :, 0] *= 2 * self.window_size[1] - 1
        relative_position_index = relative_coords.sum(-1)

        self.relative_position_index = tf.Variable(
            initial_value=tf.convert_to_tensor(relative_position_index), trainable=False
        )

    def call(self, x, mask=None):
        _, size, channels = x.shape
        head_dim = channels // self.num_heads
        x_qkv = self.qkv(x)
        x_qkv = tf.reshape(x_qkv, shape=(-1, size, 3, self.num_heads, head_dim))
        x_qkv = tf.transpose(x_qkv, perm=(2, 0, 3, 1, 4))
        q, k, v = x_qkv[0], x_qkv[1], x_qkv[2]
        q = q * self.scale
        k = tf.transpose(k, perm=(0, 1, 3, 2))
        attn = q @ k

        num_window_elements = self.window_size[0] * self.window_size[1]
        relative_position_index_flat = tf.reshape(
            self.relative_position_index, shape=(-1, )
        )
        relative_position_bias = tf.gather(
            self.relative_position_bias_table, relative_position_index_flat
        )
        relative_position_bias = tf.reshape(
            relative_position_bias, shape=(num_window_elements, num_window_elements, -1)
        )
        relative_position_bias = tf.transpose(relative_position_bias, perm=(2, 0, 1))
        attn = attn + tf.expand_dims(relative_position_bias, axis=0)

        if mask is not None:
            nW = mask.get_shape()[0]
            mask_float = tf.cast(
                tf.expand_dims(tf.expand_dims(mask, axis=1), axis=0), tf.float32
            )
            attn = (
                tf.reshape(attn, shape=(-1, nW, self.num_heads, size, size))
                + mask_float
            )
            attn = tf.reshape(attn, shape=(-1, self.num_heads, size, size))
            attn = keras.activations.softmax(attn, axis=-1)
        else:
            attn = keras.activations.softmax(attn, axis=-1)
        attn = self.dropout(attn)

        x_qkv = attn @ v
        x_qkv = tf.transpose(x_qkv, perm=(0, 2, 1, 3))
        x_qkv = tf.reshape(x_qkv, shape=(-1, size, channels))
        x_qkv = self.proj(x_qkv)
        x_qkv = self.dropout(x_qkv)
        return x_qkv
```

```
In [19]: class SwinTransformer(layers.Layer):
    def __init__(self,
                 dim,
                 num_patch,
                 num_heads,
                 window_size=7,
                 shift_size=0,
                 num_mlp=1024,
                 qkv_bias=True,
                 dropout_rate=0.0,
                 **kwargs,
                 ):
        super(SwinTransformer, self).__init__(**kwargs)

        self.dim = dim # number of input dimensions
        self.num_patch = num_patch # number of embedded patches
        self.num_heads = num_heads # number of attention heads
        self.window_size = window_size # size of window
        self.shift_size = shift_size # size of window shift
        self.num_mlp = num_mlp # number of MLP nodes

        self.norm1 = layers.LayerNormalization(epsilon=1e-5)
        self.attn = WindowAttention(
            dim,
            window_size=(self.window_size, self.window_size),
            num_heads=num_heads,
            qkv_bias=qkv_bias,
            dropout_rate=dropout_rate,
        )
        self.drop_path = DropPath(dropout_rate)
        self.norm2 = layers.LayerNormalization(epsilon=1e-5)

        self.mlp = keras.Sequential(
            [
                layers.Dense(num_mlp),
                layers.Activation(keras.activations.gelu),
                layers.Dropout(dropout_rate),
                layers.Dense(dim),
                layers.Dropout(dropout_rate),
            ]
        )

        if min(self.num_patch) < self.window_size:
            self.shift_size = 0
            self.window_size = min(self.num_patch)

    def build(self, input_shape):
        if self.shift_size == 0:
            self.attn_mask = None
        else:
            height, width = self.num_patch
            h_slices = (
                slice(0, -self.window_size),
                slice(-self.window_size, -self.shift_size),
                slice(-self.shift_size, None),
            )
            w_slices = (
                slice(0, -self.window_size),
                slice(-self.window_size, -self.shift_size),
                slice(-self.shift_size, None),
            )
            mask_array = np.zeros((1, height, width, 1))
            count = 0
            for h in h_slices:
                for w in w_slices:
                    mask_array[:, h, w, :] = count
                    count += 1
            mask_array = tf.convert_to_tensor(mask_array)

            # mask array to windows
            mask_windows = window_partition(mask_array, self.window_size)
            mask_windows = tf.reshape(
                mask_windows, shape=[-1, self.window_size * self.window_size]
            )
            attn_mask = tf.expand_dims(mask_windows, axis=1) - tf.expand_dims(
                mask_windows, axis=2
            )
            attn_mask = tf.where(attn_mask != 0, -100.0, attn_mask)
            attn_mask = tf.where(attn_mask == 0, 0.0, attn_mask)
            self.attn_mask = tf.Variable(initial_value=attn_mask, trainable=False)

    def call(self, x):
        height, width = self.num_patch
        _, num_patches_before, channels = x.shape
        x_skip = x
        x = self.norm1(x)
        x = tf.reshape(x, shape=(-1, height, width, channels))
        if self.shift_size > 0:
            shifted_x = tf.roll(
                x, shift=[-self.shift_size, -self.shift_size], axis=[1, 2]
            )
        else:
            shifted_x = x

        x_windows = window_partition(shifted_x, self.window_size)
        x_windows = tf.reshape(
            x_windows, shape=[-1, self.window_size * self.window_size, channels]
        )
        attn_windows = self.attn(x_windows, mask=self.attn_mask)

        attn_windows = tf.reshape(
            attn_windows, shape=(-1, self.window_size, self.window_size, channels)
        )
        shifted_x = window_reverse(
            attn_windows, self.window_size, height, width, channels
        )
        if self.shift_size > 0:
            x = tf.roll(
                shifted_x, shift=[self.shift_size, self.shift_size], axis=[1, 2]
            )
        else:
            x = shifted_x

        x = tf.reshape(x, shape=(-1, height * width, channels))
        x = self.drop_path(x)
        x = x_skip + x
        x_skip = x
        x = self.norm2(x)
        x = self.mlp(x)
        x = self.drop_path(x)
        x = x_skip + x
        return x
```

```
In [20]: class PatchExtract(layers.Layer):
    def __init__(self, patch_size, **kwargs):
        super(PatchExtract, self).__init__(**kwargs)
        self.patch_size_x = patch_size[0]
        self.patch_size_y = patch_size[1]

    def call(self, images):
        batch_size = tf.shape(images)[0]
        patches = tf.image.extract_patches(
            images=images,
            sizes=[1, self.patch_size_x, self.patch_size_y, 1],
            strides=(1, self.patch_size_x, self.patch_size_y, 1),
            rates=(1, 1, 1, 1),
            padding="VALID",
        )
        patch_dim = patches.shape[-1]
        patch_num = patches.shape[1]
        return tf.reshape(patches, (batch_size, patch_num * patch_num, patch_dim))

class PatchEmbedding(layers.Layer):
    def __init__(self, num_patch, embed_dim, **kwargs):
        super(PatchEmbedding, self).__init__(**kwargs)
        self.num_patch = num_patch
        self.proj = layers.Dense(embed_dim)
        self.pos_embed = layers.Embedding(input_dim=num_patch, output_dim=embed_dim)

    def call(self, patch):
        pos = tf.range(start=0, limit=self.num_patch, delta=1)
        return self.proj(patch) + self.pos_embed(pos)

class PatchMerging(tf.keras.layers.Layer):
    def __init__(self, num_patch, embed_dim):
        super(PatchMerging, self).__init__()
        self.num_patch = num_patch
        self.embed_dim = embed_dim
        self.linear_trans = layers.Dense(2 * embed_dim, use_bias=False)

    def call(self, x):
        height, width = self.num_patch
        _, _, C = x.get_shape().as_list()
        x = tf.reshape(x, shape=(-1, height, width, C))
        x0 = x[:, 0::2, 0::2, :]
        x1 = x[:, 1::2, 0::2, :]
        x2 = x[:, 0::2, 1::2, :]
        x3 = x[:, 1::2, 1::2, :]
        x = tf.concat((x0, x1, x2, x3), axis=-1)
        x = tf.reshape(x, shape=(-1, (height // 2) * (width // 2), 4 * C))
        return self.linear_trans(x)
```

```
In [21]: input = layers.Input(input_shape)
x = layers.RandomCrop(image_dimension, image_dimension)(input)
x = layers.RandomFlip("horizontal")(x)
x = PatchExtract(patch_size)(x)
x = PatchEmbedding(num_patch_x * num_patch_y, embed_dim)(x)
x = SwinTransformer(
    dim=embed_dim,
    num_patch=(num_patch_x, num_patch_y),
    num_heads=num_heads,
    window_size=window_size,
    shift_size=0,
    num_mlp=num_mlp,
    qkv_bias=qkv_bias,
    dropout_rate=dropout_rate,
)(x)
x = SwinTransformer(
    dim=embed_dim,
    num_patch=(num_patch_x, num_patch_y),
    num_heads=num_heads,
    window_size=window_size,
    shift_size=shift_size,
    num_mlp=num_mlp,
    qkv_bias=qkv_bias,
    dropout_rate=dropout_rate,
)(x)
x = PatchMerging((num_patch_x, num_patch_y), embed_dim=embed_dim)(x)
x = layers.GlobalAveragePooling1D()(x)
output = layers.Dense(num_classes, activation="softmax")(x)
```

WARNING:tensorflow:Using a while\_loop for converting RngReadAndSkip cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting StatelessRandomUniformIntV2 cause there is no registered converter for this op.

```
In [22]: model1 = keras.Model(input, output)
model1.compile(loss = 'categorical_crossentropy', optimizer='adam', metrics=["accuracy",f1_m,precision_m, recall_m])
model1.summary()
```

```
Model: "model"
-----  

Layer (type)          Output Shape         Param #
-----  

input_1 (InputLayer)   [(None, 256, 256, 3)]  0  

random_crop (RandomCrop) (None, 256, 256, 3)  0  

random_flip (RandomFlip) (None, 256, 256, 3)  0  

patch_extract (PatchExtract (None, 16384, 12)  0  

)  

patch_embedding (PatchEmbed (None, 16384, 64)  1049408  

ding)  

swin_transformer (SwinTrans (None, 16384, 64)  50072  

former)  

swin_transformer_1 (SwinTra (None, 16384, 64)  115608  

nsformer)  

patch_merging (PatchMergin (None, 4096, 128)  32768  

)  

global_average_pooling1d (G (None, 128)  0  

lobalAveragePooling1D)  

dense_10 (Dense)      (None, 4)           516  

-----  

Total params: 1,248,372  

Trainable params: 1,182,804  

Non-trainable params: 65,568
```

```
In [24]: hist1 = model1.fit(train_set, validation_data=test_set, epochs=50, steps_per_epoch=len(train_set), validation_steps=len(test_set))#, callbacks=[learning_rate_reduction, early_stop]
```

Epoch 1/50  
 WARNING:tensorflow:Using a while\_loop for converting RngReadAndSkip cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting RngReadAndSkip cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.  
 WARNING:tensorflow:Using a while\_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.  
 2856/2856 [=====] - 134s 44ms/step - loss: 0.9785 - accuracy: 0.5942 - f1\_m: 0.5037 - precision\_m: 0.5921 - recall\_m: 0.4596 - val\_loss: 0.8797 - val\_accuracy: 0.6026 - val\_f1\_m: 0.5401 - val\_precision\_m: 0.6250 - val\_recall\_m: 0.4977  
 Epoch 2/50  
 2856/2856 [=====] - 108s 38ms/step - loss: 0.6742 - accuracy: 0.7384 - f1\_m: 0.7082 - precision\_m: 0.7719 - recall\_m: 0.6763 - val\_loss: 0.8228 - val\_accuracy: 0.7155 - val\_f1\_m: 0.6880 - val\_precision\_m: 0.7363 - val\_recall\_m: 0.6639  
 Epoch 3/50  
 2856/2856 [=====] - 109s 38ms/step - loss: 0.6203 - accuracy: 0.7647 - f1\_m: 0.7382 - precision\_m: 0.7971 - recall\_m: 0.7087 - val\_loss: 0.7658 - val\_accuracy: 0.7033 - val\_f1\_m: 0.6641 - val\_precision\_m: 0.7134 - val\_recall\_m: 0.6395  
 Epoch 4/50  
 2856/2856 [=====] - 108s 38ms/step - loss: 0.5970 - accuracy: 0.7747 - f1\_m: 0.7545 - precision\_m: 0.8104 - recall\_m: 0.7265 - val\_loss: 0.6724 - val\_accuracy: 0.7265 - val\_f1\_m: 0.7545 - val\_precision\_m: 0.8104 - val\_recall\_m: 0.7265

```
In [66]: dl_acc = hist1.history["val_accuracy"][49]
dl_prec = hist1.history["val_precision_m"][49]
dl_rec = hist1.history["val_recall_m"][49]
dl_f1 = hist1.history["val_f1_m"][49]
```



```
storeResults('VisionTransformer - SWIN',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [25]: import matplotlib.pyplot as plt
```

```
x=hist1
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## CCT

```
In [26]: # Model
from keras.layers import add
from keras.layers import Dense
from keras.layers import Layer
from keras.layers import Input
from keras.layers import Conv2D
from keras.layers import Dropout
from keras.layers import Embedding
from keras.layers import MaxPooling2D as MP
from keras.layers import ZeroPadding2D as ZP
from keras.layers import LayerNormalization as LN
from keras.layers import MultiHeadAttention as MHA
from keras.models import Sequential, Model

# Optimizer
from tensorflow_addons.optimizers import AdamW

# Model Visualization
from tensorflow.keras.utils import plot_model
```

```
In [27]: Learning_Rate = 1e-3
Weight_Decay = 1e-4
FILTERS = [64, 256]
Drop_Rates = [0.0, 0.1]
Dims = 256
Hidden_units = [Dims, Dims]
IMAGE_SIZE = 256
```

```
In [28]: def show_images(data, class_names, model=None, GRID=[5,6], SIZE=(30,25)):

    # Plotting Configuration
    n_rows, n_cols = GRID
    n_images = n_rows * n_cols
    plt.figure(figsize=SIZE)

    # iterate through the data
    i=1
    for images, labels in iter(data):

        # Select data at random
        id = np.random.randint(len(images))
        image, label = tf.expand_dims(images[id], axis=0), class_names[int(labels[id])]

        # Make Prediction
        if model is not None :
            prediction = model.predict(image)[0]
            score = np.round(np.max(prediction),2)
            pred = class_names[np.argmax(prediction)]

            title = "True : {}\nPred : {}\n Score : {:.2f}".format(label, pred, score)

        else:
            title = label
            cls()

        # Plot Image
        plt.subplot(n_rows, n_cols, i)
        plt.imshow(image[0])
        plt.axis('off')
        plt.title(title)

        # Break Loop
        i+=1
        if i>n_images: break

    # Show Final Figure
    plt.show()
```

```
In [29]: # get the Class Names
import os
class_names = sorted(os.listdir(train_path))
n_classes = len(class_names)

# Show
print("No. of Classes : {} \nClass Names : {}".format(n_classes, class_names))

No. of Classes : 4
Class Names : ['glioma', 'meningioma', 'notumor', 'pituitary']
```

```
In [30]: class ConvTokenizer(Layer):

    def __init__(self, filters, **kwargs):
        super(ConvTokenizer, self).__init__(**kwargs)

        self.net = Sequential([
            Conv2D(filters[0], kernel_size=3, strides=1, padding='valid', activation='relu', kernel_initializer='he_normal', use_bias=False, name="TokenizerConv1"),
            ZP(padding=1, name="ZP1"),
            MP(pool_size=3, strides=2, padding='same', name="MP1"),
            Conv2D(filters[1], kernel_size=3, strides=1, padding='valid', activation='relu', kernel_initializer='he_normal', use_bias=False, name="TokenizerConv2"),
            ZP(padding=1, name="ZP2"),
            MP(pool_size=3, strides=2, padding='same', name="MP2"),
        ], name="ConvTokenizer")

        self.filters = filters

    def call(self, X):
        patches = self.net(X)
        patches = tf.reshape(patches, shape=(-1, patches.shape[1]*patches.shape[2], patches.shape[-1]))
        return patches

    def PE(self, IMAGE_SIZE):
        dummy_inputs = tf.ones(shape=(1, IMAGE_SIZE, IMAGE_SIZE, 3))
        dummy_outputs = self.call(dummy_inputs)

        seq_len = tf.shape(dummy_outputs)[1]
        proj_dim = tf.shape(dummy_outputs)[-1]
        EL = Embedding(input_dim=seq_len, output_dim=proj_dim, name="PositionalEmbedding")
        return EL, seq_len

    def get_config(self):
        base_config = super().get_config()
        return {**base_config, "filters":self.filters}
```

```
In [31]: def SeqPool(x):

    rep = LN(epsilon=1e-5, name="LayerNormSeqPool")(x)
    y = Dense(1, name="LinearTransformation")(rep)

    attention_weights = tf.nn.softmax(y, axis=-1)
    weighted_rep = tf.squeeze(tf.matmul(attention_weights, rep, transpose_a=True), axis=-2)
    return weighted_rep
```

```
In [32]: class MLPBlock(Layer):
```

```
    def __init__(self, units, rate, **kwargs):
        super(MLPBlock, self).__init__(**kwargs)

        self.units = units
        self.rate = rate

        self.net = Sequential(name='MLPBlock')
        for unit in units:
            self.net.add(Dense(unit, activation='relu', kernel_initializer='he_normal'))
        self.net.add(Dropout(rate))

    def call(self, X):
        return self.net(X)

    def get_config(self):
        base_config = super().get_config()
        return {
            **base_config,
            "units":self.units,
            "rate":self.rate
        }
```

```
In [33]: class StochasticDepthRegularization(Layer):
```

```
    def __init__(self, drop_rate, **kwargs):
        super(StochasticDepthRegularization, self).__init__(**kwargs)

        self.dr = drop_rate

    def call(self, x, training=None):
        if training:
            keep_rate = 1 - self.dr
            shape = (tf.shape(x)[0],) + (1,) * (tf.shape(x).shape[0] - 1)
            random_tensor = keep_rate + tf.random.uniform(shape, 0, 1)
            random_tensor = tf.floor(random_tensor)
            return (x / keep_rate) * random_tensor
        return x
```

```
In [34]: def TransformerBlock(x, rates=Drop_Rates, L=2, proj_dim=Dims):
```

```
    for i in range(L):
        y = LN(epsilon=1e-5, name="LayerNormTB_Lower_{}".format(i+1))(x)
        y = MHA(num_heads=2, key_dim=Dims, dropout=0.1, name="MHA_{}".format(i+1))(y, y)
        y = StochasticDepthRegularization(drop_rate=rates[i], name="SDR_Lower_{}".format(i+1))(y)

        skip = add([y, x], name="SkipConnection_Lower_{}".format(i+1))

        y = LN(epsilon=1e-5, name="LayerNormTB_Top_{}".format(i+1))(skip)
        y = MLPBlock(Hidden_units, 0.1, name="MLPBlock_TB_{}".format(i+1))(y)
        y = StochasticDepthRegularization(drop_rate=rates[i], name="SDR_Top_{}".format(i+1))(y)

        x = add([y, skip], name="SkipConnection_Top_{}".format(i+1))
    return x
```

```
In [35]: # Input Layer
InputL = Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 3), name="InputLayer")

# Tokenizer
Tokenizer = ConvTokenizer(filters=FILTERS, name="ConvTokenizer")
tokens = Tokenizer(InputL)

# Positional Embeddings
EL, seq_len = Tokenizer.PE(IMAGE_SIZE)
positions = tf.range(start=0, limit=seq_len, delta=1)
embedding = EL(positions)
tokens += embedding

# Transformer Block
encodings = TransformerBlock(tokens)

# Sequence Pooling
pooled = SeqPool(encodings)

# Classifier
OutputL = Dense(n_classes, activation='softmax', name="OutputLayer")(pooled)

# Model
model2 = Model(InputL, OutputL, name="CCT")
model2.summary()
```

Model: "CCT"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
InputLayer (InputLayer)	[None, 256, 256, 3]	0	[]
ConvTokenizer (ConvTokenizer)	(None, 4096, 256)	149184	['InputLayer[0][0]']
tf.__operators__.add (TFOpLamb da)	(None, 4096, 256)	0	['ConvTokenizer[0][0]']
LayerNormTB_Lower_1 (LayerNorm alization)	(None, 4096, 256)	512	['tf.__operators__.add[0][0]']
MHA_1 (MultiHeadAttention)	(None, 4096, 256)	526080	['LayerNormTB_Lower_1[0][0]', 'LayerNormTB_Lower_1[0][0]']
SDR_Lower_1 (StochasticDepthRe gularization)	(None, 4096, 256)	0	['MHA_1[0][0]']
SkipConnection_Lower_1 (Add)	(None, 4096, 256)	0	['SDR_Lower_1[0][0]', 'tf.__operators__.add[0][0]']
LayerNormTB_Top_1 (LayerNormal ization)	(None, 4096, 256)	512	['SkipConnection_Lower_1[0][0]']
MLPBlock_TB_1 (MLPBlock)	(None, 4096, 256)	131584	['LayerNormTB_Top_1[0][0]']
SDR_Top_1 (StochasticDepthRegu larization)	(None, 4096, 256)	0	['MLPBlock_TB_1[0][0]']
SkipConnection_Top_1 (Add)	(None, 4096, 256)	0	['SDR_Top_1[0][0]', 'SkipConnection_Lower_1[0][0]']
LayerNormTB_Lower_2 (LayerNorm alization)	(None, 4096, 256)	512	['SkipConnection_Top_1[0][0]']
MHA_2 (MultiHeadAttention)	(None, 4096, 256)	526080	['LayerNormTB_Lower_2[0][0]', 'LayerNormTB_Lower_2[0][0]']
SDR_Lower_2 (StochasticDepthRe gularization)	(None, 4096, 256)	0	['MHA_2[0][0]']
SkipConnection_Lower_2 (Add)	(None, 4096, 256)	0	['SDR_Lower_2[0][0]', 'SkipConnection_Top_1[0][0]']
LayerNormTB_Top_2 (LayerNormal ization)	(None, 4096, 256)	512	['SkipConnection_Lower_2[0][0]']
MLPBlock_TB_2 (MLPBlock)	(None, 4096, 256)	131584	['LayerNormTB_Top_2[0][0]']
SDR_Top_2 (StochasticDepthRegu larization)	(None, 4096, 256)	0	['MLPBlock_TB_2[0][0]']
SkipConnection_Top_2 (Add)	(None, 4096, 256)	0	['SDR_Top_2[0][0]', 'SkipConnection_Lower_2[0][0]']
LayerNormSeqPool (LayerNormali zation)	(None, 4096, 256)	512	['SkipConnection_Top_2[0][0]']
LinearTransformation (Dense)	(None, 4096, 1)	257	['LayerNormSeqPool[0][0]']
tf.nn.softmax (TFOpLambda)	(None, 4096, 1)	0	['LinearTransformation[0][0]']
tf.linalg.matmul (TFOpLambda)	(None, 1, 256)	0	['tf.nn.softmax[0][0]', 'LayerNormSeqPool[0][0]']
tf.compat.v1.squeeze (TFOpLamb da)	(None, 256)	0	['tf.linalg.matmul[0][0]']
OutputLayer (Dense)	(None, 4)	1028	['tf.compat.v1.squeeze[0][0]']
<hr/>			

Total params: 1,468,357  
Trainable params: 1,468,357  
Non-trainable params: 0

```
In [36]: opt = AdamW(learning_rate=Learning_Rate, weight_decay=Weight_Decay)
model2.compile(loss = 'categorical_crossentropy', optimizer=opt, metrics=["accuracy", f1_m, precision_m, recall_m])
```

```
In [37]: hist2 = model2.fit(train_set, validation_data=test_set, epochs=50, steps_per_epoch=len(train_set), validation_steps=len(test_set))#, callbacks=[learning_rate_reduction, early_stop]
```

Epoch 1/50  
2856/2856 [=====] - 127s 41ms/step - loss: 414.0080 - accuracy: 0.4597 - f1\_m: 0.4597 - precision\_m: 0.4597 - recall\_m: 0.4597 - val\_loss: 145.4632 - val\_accuracy: 0.3638 - val\_f1\_m: 0.3643 - val\_precision\_m: 0.3643 - val\_recall\_m: 0.3643  
Epoch 2/50  
2856/2856 [=====] - 117s 41ms/step - loss: 51.3703 - accuracy: 0.5506 - f1\_m: 0.5507 - precision\_m: 0.5508 - recall\_m: 0.5506 - val\_loss: 10.8244 - val\_accuracy: 0.6506 - val\_f1\_m: 0.6512 - val\_precision\_m: 0.6532 - val\_recall\_m: 0.6502  
Epoch 3/50  
2856/2856 [=====] - 117s 41ms/step - loss: 11.0667 - accuracy: 0.5807 - f1\_m: 0.5809 - precision\_m: 0.5819 - recall\_m: 0.5804 - val\_loss: 5.5326 - val\_accuracy: 0.4882 - val\_f1\_m: 0.4886 - val\_precision\_m: 0.4886 - val\_recall\_m: 0.4886  
Epoch 4/50  
2856/2856 [=====] - 117s 41ms/step - loss: 1.5594 - accuracy: 0.6122 - f1\_m: 0.5961 - precision\_m: 0.6285 - recall\_m: 0.5798 - val\_loss: 2.2409 - val\_accuracy: 0.4218 - val\_f1\_m: 0.4009 - val\_precision\_m: 0.4284 - val\_recall\_m: 0.3872  
Epoch 5/50  
2856/2856 [=====] - 117s 41ms/step - loss: 1.3099 - accuracy: 0.5550 - f1\_m: 0.5244 - precision\_m: 0.5772 - recall\_m: 0.4981 - val\_loss: 0.8488 - val\_accuracy: 0.7010 - val\_f1\_m: 0.5526 - val\_precision\_m: 0.6806 - val\_recall\_m: 0.4886  
Epoch 6/50  
2856/2856 [=====] - 117s 41ms/step - loss: 1.3225 - accuracy: 0.5572 - f1\_m: 0.5275 - precision\_m: 0.5776 - recall\_m: 0.5025 - val\_loss: 1.1885 - val\_accuracy: 0.5995 - val\_f1\_m: 0.5498 - val\_precision\_m: 0.6296 - val\_recall\_m: 0.5099  
Epoch 7/50  
.../50



```
In [67]: dl_acc = hist2.history["val_accuracy"][49]
dl_prec = hist2.history["val_precision_m"][49]
dl_rec = hist2.history["val_recall_m"][49]
dl_f1 = hist2.history["val_f1_m"][49]
```

```
storeResults('VisionTransformer - CCT',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [38]: x=hist2
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## EANet

```
In [39]: import keras
from keras import layers
```

```
In [40]: learning_rate = 0.001
weight_decay = 0.0001
batch_size = 128
num_epochs = 10
```

```
In [41]: image_size = 256
auto = tf.data.AUTOTUNE

augmentation_layers = [
    keras.layers.RandomCrop(256, 256),
    keras.layers.RandomFlip("horizontal"),
]
```

```
In [42]: def activation_block(x):
    x = layers.Activation("gelu")(x)
    return layers.BatchNormalization()(x)

def conv_stem(x, filters: int, patch_size: int):
    x = layers.Conv2D(filters, kernel_size=patch_size, strides=patch_size)(x)
    return activation_block(x)

def conv_mixer_block(x, filters: int, kernel_size: int):
    # Depthwise convolution.
    x0 = x
    x = layers.DepthwiseConv2D(kernel_size=kernel_size, padding="same")(x)
    x = layers.Add()([activation_block(x), x0]) # Residual.

    # Pointwise convolution.
    x = layers.Conv2D(filters, kernel_size=1)(x)
    x = activation_block(x)

    return x

def eanet(
    image_size=32, filters=768, depth=8, kernel_size=5, patch_size=2, num_classes=7
):
    inputs = keras.Input((image_size, image_size, 3))
    x = layers.Rescaling(scale=1.0 / 255)(inputs)

    # Extract patch embeddings.
    x = conv_stem(x, filters, patch_size)

    # ConvMixer blocks.
    for _ in range(depth):
        x = conv_mixer_block(x, filters, kernel_size)

    # Classification block.
    x = layers.GlobalAvgPool2D()(x)
    outputs = layers.Dense(4, activation="softmax")(x)

    return keras.Model(inputs, outputs)
```

```
In [43]: model2 = eanet()
```

```
In [44]: model2.compile(loss = 'categorical_crossentropy', optimizer='adam', metrics=["accuracy",f1_m,precision_m, recall_m])
model2.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_2 (InputLayer)	[None, 32, 32, 3]	0	[]
rescaling (Rescaling)	(None, 32, 32, 3)	0	['input_2[0][0]']
conv2d (Conv2D)	(None, 16, 16, 768)	9984	['rescaling[0][0]']
activation_2 (Activation)	(None, 16, 16, 768)	0	['conv2d[0][0]']
batch_normalization (BatchNorm alization)	(None, 16, 16, 768)	3072	['activation_2[0][0]']
depthwise_conv2d (DepthwiseCon v2D)	(None, 16, 16, 768)	19968	['batch_normalization[0][0]']
activation_3 (Activation)	(None, 16, 16, 768)	0	['depthwise_conv2d[0][0]']

```
In [45]: hist2a = model2.fit(train_set, validation_data=test_set, epochs=50, steps_per_epoch=len(train_set), validation_steps=len(test_set))#, callbacks=[learning_rate_reduction, early_stop]
```

Epoch 1/50  
2856/2856 [=====] - 496s 172ms/step - loss: 1.1885 - accuracy: 0.4807 - f1\_m: 0.2741 - precision\_m: 0.3528 - recall\_m: 0.2348 - val\_loss: 2.0107 - val\_accuracy: 0.2708 - val\_f1\_m: 0.2612 - val\_precision\_m: 0.2851 - val\_recall\_m: 0.2492  
Epoch 2/50  
2856/2856 [=====] - 492s 172ms/step - loss: 1.0829 - accuracy: 0.5513 - f1\_m: 0.3846 - precision\_m: 0.4795 - recall\_m: 0.3372 - val\_loss: 12.5834 - val\_accuracy: 0.2403 - val\_f1\_m: 0.2411 - val\_precision\_m: 0.2416 - val\_recall\_m: 0.2409  
Epoch 3/50  
2856/2856 [=====] - 492s 172ms/step - loss: 0.9928 - accuracy: 0.6033 - f1\_m: 0.4950 - precision\_m: 0.5917 - recall\_m: 0.4466 - val\_loss: 47.5491 - val\_accuracy: 0.2288 - val\_f1\_m: 0.2287 - val\_precision\_m: 0.2287 - val\_recall\_m: 0.2287  
Epoch 4/50  
2856/2856 [=====] - 492s 172ms/step - loss: 0.9022 - accuracy: 0.6576 - f1\_m: 0.5778 - precision\_m: 0.6637 - recall\_m: 0.5348 - val\_loss: 110.7900 - val\_accuracy: 0.2288 - val\_f1\_m: 0.2294 - val\_precision\_m: 0.2294 - val\_recall\_m: 0.2294  
Epoch 5/50  
2856/2856 [=====] - 492s 172ms/step - loss: 0.8473 - accuracy: 0.6919 - f1\_m: 0.6257 - precision\_m: 0.7034 - recall\_m: 0.5868 - val\_loss: 8.7278 - val\_accuracy: 0.3661 - val\_f1\_m: 0.3651 - val\_precision\_m: 0.3666 - val\_recall\_m: 0.3643  
Epoch 6/50  
2856/2856 [=====] - 492s 172ms/step - loss: 0.8072 - accuracy: 0.7132 - f1\_m: 0.6531 - precision\_m: 0.7180 - recall\_m: 0.6206 - val\_loss: 10.2742 - val\_accuracy: 0.2517 - val\_f1\_m: 0.2508 - val\_precision\_m: 0.2508 - val\_recall\_m: 0.2508  
Epoch 7/50  
2856/2856 [=====] - 492s 172ms/step - loss: 0.7740 - accuracy: 0.7414 - f1\_m: 0.6926 - precision\_m: 0.7411 - recall\_m: 0.6522 - val\_loss: 110.2052 - val\_accuracy: 0.3000 - val\_f1\_m: 0.2994 - val\_precision\_m: 0.3000 - val\_recall\_m: 0.2994



```
In [68]: dl_acc = hist2a.history["val_accuracy"][49]
dl_prec = hist2a.history["val_precision_m"][49]
dl_rec = hist2a.history["val_recall_m"][49]
dl_f1 = hist2a.history["val_f1_m"][49]
```

```
storeResults('VisionTransformer - EANet',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [46]: x=hist2a
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=16)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## TL Models

```
In [69]: from tensorflow.keras.layers import Dense, Flatten, Input, Lambda
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.applications import VGG16
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img
from glob import glob
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [24]: from tensorflow.keras.layers import *
from tensorflow.keras.losses import *
from tensorflow.keras.models import *
from tensorflow.keras.optimizers import *
from sklearn.metrics import confusion_matrix
```

```
In [25]: # Resizing all the images to (224,224)
IMAGE_SIZE = [256,256]

train_path = 'Brain_Tumor_MRI_Image_Dataset//Training'
test_path = 'Brain_Tumor_MRI_Image_Dataset//Testing'
```

```
In [26]: # Scaling all the images between 0 to 1

train_datagen = ImageDataGenerator(rescale = 1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=False)

# Performing only scaling on the test dataset

test_datagen = ImageDataGenerator(rescale=1./255)
```

```
In [27]: train_set = train_datagen.flow_from_directory(train_path,
                                                 target_size=(256,256),
                                                 batch_size=2,
                                                 class_mode = 'categorical')

test_set = test_datagen.flow_from_directory(test_path,
                                            target_size=(256,256),
                                            batch_size=2,
                                            class_mode='categorical')
```

Found 5712 images belonging to 4 classes.  
Found 1311 images belonging to 4 classes.

## VGG16

```
In [28]: base_model = VGG16(input_shape = (256,256,3), weights=None, include_top=True)

x1= Flatten()(base_model.output)
prediction1 = Dense(4, activation='softmax')(x1)

model3 = Model(inputs = base_model.inputs, outputs = prediction1)
model3.summary()
model3.compile(loss = 'categorical_crossentropy', optimizer=Adam(learning_rate=0.0001), metrics=["accuracy",f1_m,precision_m, recall_m])
<

Model: "model_2"

```

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[None, 256, 256, 3]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
flatten (Flatten)	(None, 32768)	0
fc1 (Dense)	(None, 4096)	134221824
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
flatten_2 (Flatten)	(None, 1000)	0
dense_2 (Dense)	(None, 4)	4004

```
=====
Total params: 169,818,828
Trainable params: 169,818,828
Non-trainable params: 0

```

```
In [29]: hist3 = model3.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])
<

Epoch 1/50
2856/2856 [=====] - 82s 28ms/step - loss: 1.3843 - accuracy: 0.2771 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
21 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 1.0000e-04
Epoch 2/50
2856/2856 [=====] - 79s 28ms/step - loss: 1.3837 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
14 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 1.0000e-04
Epoch 3/50
2856/2856 [=====] - 79s 28ms/step - loss: 1.3836 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
10 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 1.0000e-04
Epoch 4/50
2856/2856 [=====] - ETA: 0s - loss: 1.3835 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 4: ReduceLROnPlateau reducing learning rate to 2.9999999242136255e-05.
2856/2856 [=====] - 83s 29ms/step - loss: 1.3835 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
08 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 1.0000e-04
Epoch 5/50
2856/2856 [=====] - 83s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
08 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 3.0000e-05
Epoch 6/50
2856/2856 [=====] - 84s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
07 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 3.0000e-05
Epoch 7/50
2855/2856 [=====] - ETA: 0s - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 7: ReduceLROnPlateau reducing learning rate to 8.99999772640877e-06.
2856/2856 [=====] - 86s 30ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
07 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 3.0000e-05
Epoch 8/50
2856/2856 [=====] - 83s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
07 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-06
Epoch 9/50
2856/2856 [=====] - 83s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
06 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-06
Epoch 10/50
2856/2856 [=====] - ETA: 0s - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 10: ReduceLROnPlateau reducing learning rate to 2.699998226528985e-06.
2856/2856 [=====] - 83s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
06 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-06
Epoch 11/50
2855/2856 [=====] - ETA: 0s - loss: 1.3834 - accuracy: 0.2793 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00Restoring model weights
from the end of the best epoch: 1.
2856/2856 [=====] - 84s 29ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
06 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 2.7000e-06
Epoch 11: early stopping

```

```
In [70]: predictions = model3.predict(test_set)
y_pred = np.argmax(predictions, axis=1)
y_true = test_set.classes

dl_acc = hist3.history["val_accuracy"][10]
dl_prec = precision_score(y_true, y_pred, average='weighted')
dl_rec = recall_score(y_true, y_pred, average='weighted')
dl_f1 = f1_score(y_true, y_pred, average='weighted')

storeResults('TL - VGG16', dl_acc, dl_prec, dl_rec, dl_f1)



656/656 [=====] - 4s 6ms/step
```

```
In [30]: model3.save('models/vgg16.h5')
```

```
In [31]: x=hist3
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## VGG19

```
In [32]: base_model = VGG19(input_shape = (256,256,3), weights=None, include_top=True)

x1= Flatten()(base_model.output)
prediction1 = Dense(4, activation='softmax')(x1)

model4 = Model(inputs = base_model.inputs, outputs = prediction1)
model4.summary()
model4.compile(loss = 'categorical_crossentropy', optimizer='sgd', metrics=["accuracy",f1_m,precision_m, recall_m])

Model: "model_3"

```

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[None, 256, 256, 3]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv4 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv4 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv4 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
flatten (Flatten)	(None, 32768)	0
fc1 (Dense)	(None, 4096)	134221824
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
flatten_3 (Flatten)	(None, 1000)	0
dense_3 (Dense)	(None, 4)	4004

```
Total params: 175,128,524
Trainable params: 175,128,524
Non-trainable params: 0
```

```
In [33]: hist4 = model4.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])

Epoch 1/50
2856/2856 [=====] - 79s 27ms/step - loss: 1.3850 - accuracy: 0.2771 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
31 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 77s 27ms/step - loss: 1.3847 - accuracy: 0.2731 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.37
96 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 77s 27ms/step - loss: 1.3849 - accuracy: 0.2777 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.37
93 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 4/50
2855/2856 [=====] - ETA: 0s - loss: 1.3847 - accuracy: 0.2764 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 4: ReduceLROnPlateau reducing learning rate to 0.002999999329447745.
2856/2856 [=====] - 77s 27ms/step - loss: 1.3847 - accuracy: 0.2763 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
11 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 77s 27ms/step - loss: 1.3839 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
06 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0030
Epoch 6/50
2856/2856 [=====] - 77s 27ms/step - loss: 1.3837 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
15 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0030
Epoch 7/50
2855/2856 [=====] - ETA: 0s - loss: 1.3839 - accuracy: 0.2793 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 7: ReduceLROnPlateau reducing learning rate to 0.000900000078231095.
2856/2856 [=====] - 77s 27ms/step - loss: 1.3839 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
02 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0030
Epoch 8/50
2856/2856 [=====] - 78s 27ms/step - loss: 1.3835 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
01 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-04
Epoch 9/50
2856/2856 [=====] - 77s 27ms/step - loss: 1.3835 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
03 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-04
Epoch 10/50
2855/2856 [=====] - ETA: 0s - loss: 1.3835 - accuracy: 0.2793 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00
Epoch 10: ReduceLROnPlateau reducing learning rate to 0.0002699999536201356.
2856/2856 [=====] - 77s 27ms/step - loss: 1.3835 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
04 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 9.0000e-04
Epoch 11/50
2856/2856 [=====] - ETA: 0s - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00Restoring model weights
from the end of the best epoch: 1.
2856/2856 [=====] - 77s 27ms/step - loss: 1.3834 - accuracy: 0.2792 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.38
04 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 2.7000e-04
Epoch 11: early stopping
```

```
In [71]: predictions = model4.predict(test_set)
y_pred = np.argmax(predictions, axis=1)
y_true = test_set.classes

dl_acc = hist4.history["val_accuracy"][10]
dl_prec = precision_score(y_true, y_pred, average='weighted')
dl_rec = recall_score(y_true, y_pred, average='weighted')
dl_f1 = f1_score(y_true, y_pred, average='weighted')



storeResults('TL - VGG19', dl_acc, dl_prec, dl_rec, dl_f1)

656/656 [=====] - 4s 6ms/step
```

```
In [34]: x=hist4
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## InceptionV3

```
In [35]: base_model = InceptionV3(input_shape = (256,256,3), weights=None, include_top=True)

x1= Flatten()(base_model.output)
prediction1 = Dense(4, activation='softmax')(x1)

model5 = Model(inputs = base_model.inputs, outputs = prediction1)
model5.summary()
model5.compile(loss = 'categorical_crossentropy', optimizer='sgd', metrics=["accuracy",f1_m,precision_m, recall_m])

Model: "model_4"
```

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_5 (InputLayer)	[(None, 256, 256, 3 0 )]	[ ]	
conv2d (Conv2D)	(None, 127, 127, 32 864 )	[ 'input_5[0][0]' ]	
batch_normalization (BatchNorm alization)	(None, 127, 127, 32 96 )	[ 'conv2d[0][0]' ]	
activation (Activation)	(None, 127, 127, 32 0 )	[ 'batch_normalization[0][0]' ]	
conv2d_1 (Conv2D)	(None, 125, 125, 32 9216 )	[ 'activation[0][0]' ]	
<hr/>			

```
In [36]: hist5 = model5.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])

Epoch 1/50
2856/2856 [=====] - 160s 54ms/step - loss: 1.3848 - accuracy: 0.2747 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3866 - val_accuracy: 0.2365 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 153s 53ms/step - loss: 1.3841 - accuracy: 0.2742 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3756 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 153s 53ms/step - loss: 1.3838 - accuracy: 0.2770 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3704 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 152s 53ms/step - loss: 1.3832 - accuracy: 0.2785 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3591 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 153s 54ms/step - loss: 1.3376 - accuracy: 0.3477 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.1976 - val_accuracy: 0.4760 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - 152s 53ms/step - loss: 1.2131 - accuracy: 0.4354 - f1_m: 0.1975 - precision_m: 0.2962 - recall_m: 0.1481 - val_loss: 1.1147 - val_accuracy: 0.4760 - val_f1_m: 0.3064 - val_precision_m: 0.4207 - val_recall_m: 0.2492 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 152s 53ms/step - loss: 1.2131 - accuracy: 0.4354 - f1_m: 0.1975 - precision_m: 0.2962 - recall_m: 0.1481 - val_loss: 1.1147 - val_accuracy: 0.4760 - val_f1_m: 0.3064 - val_precision_m: 0.4207 - val_recall_m: 0.2492 - lr: 0.0100
```

```
In [72]: dl_acc = hist5.history["val_accuracy"][24]
dl_prec = hist5.history["val_precision_m"][24]
dl_rec = hist5.history["val_recall_m"][24]
dl_f1 = hist5.history["val_f1_m"][24]



storeResults('TL - InceptionV3',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [37]: model5.save('models/inceptionv3.h5')
```

```
In [38]: x=hist5
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## ResNet50

```
In [39]: base_model = ResNet50(input_shape = (256,256,3), weights=None, include_top=True)

x1= Flatten()(base_model.output)
prediction1 = Dense(4, activation='softmax')(x1)
model6 = Model(inputs = base_model.inputs, outputs = prediction1)
model6.summary()
model6.compile(loss = 'categorical_crossentropy', optimizer='sgd', metrics=["accuracy",f1_m,precision_m, recall_m])
```

Model: "model\_5"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_6 (InputLayer)	[(None, 256, 256, 3 0 )]		[]
conv1_pad (ZeroPadding2D)	(None, 262, 262, 3 0 )		['input_6[0][0]']
conv1_conv (Conv2D)	(None, 128, 128, 64 9472 )		['conv1_pad[0][0]']
conv1_bn (BatchNormalization)	(None, 128, 128, 64 256 )		['conv1_conv[0][0]']
conv1_relu (Activation)	(None, 128, 128, 64 0 )		['conv1_bn[0][0]']
pool1_pad (ZeroPadding2D)	(None, 130, 130, 64 0 )		['conv1_relu[0][0]']

```
In [40]: hist6 = model6.fit(train_set, epochs=50, validation_data=test_set,steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])
```

```
Epoch 1/50
2856/2856 [=====] - 106s 35ms/step - loss: 1.3359 - accuracy: 0.3398 - f1_m: 0.0082 - precision_m: 0.0123 - recall_m: 0.0061 - val_loss: 1.1903 - val_accuracy: 0.4561 - val_f1_m: 0.2734 - val_precision_m: 0.3780 - val_recall_m: 0.2210 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 101s 35ms/step - loss: 1.2287 - accuracy: 0.4048 - f1_m: 0.2187 - precision_m: 0.3281 - recall_m: 0.1640 - val_loss: 1.1176 - val_accuracy: 0.4805 - val_f1_m: 0.3148 - val_precision_m: 0.4306 - val_recall_m: 0.2569 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 100s 35ms/step - loss: 1.1446 - accuracy: 0.4403 - f1_m: 0.2421 - precision_m: 0.3631 - recall_m: 0.1815 - val_loss: 1.2509 - val_accuracy: 0.4249 - val_f1_m: 0.3468 - val_precision_m: 0.4505 - val_recall_m: 0.2950 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 100s 35ms/step - loss: 1.1167 - accuracy: 0.4466 - f1_m: 0.2416 - precision_m: 0.3624 - recall_m: 0.1812 - val_loss: 1.1019 - val_accuracy: 0.4729 - val_f1_m: 0.2861 - val_precision_m: 0.3933 - val_recall_m: 0.2325 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 100s 35ms/step - loss: 1.0965 - accuracy: 0.4527 - f1_m: 0.2465 - precision_m: 0.3697 - recall_m: 0.1849 - val_loss: 1.0875 - val_accuracy: 0.4966 - val_f1_m: 0.3516 - val_precision_m: 0.4695 - val_recall_m: 0.2927 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - 100s 35ms/step - loss: 1.0955 - accuracy: 0.4578 - f1_m: 0.2558 - precision_m: 0.3838 - recall_m: 0.1919 - val_loss: 1.0630 - val_accuracy: 0.4882 - val_f1_m: 0.3013 - val_precision_m: 0.4192 - val_recall_m: 0.2424 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 100s 35ms/step - loss: 1.0955 - accuracy: 0.4578 - f1_m: 0.2558 - precision_m: 0.3838 - recall_m: 0.1919 - val_loss: 1.0630 - val_accuracy: 0.4882 - val_f1_m: 0.3013 - val_precision_m: 0.4192 - val_recall_m: 0.2424 - lr: 0.0100
```



```
In [73]: dl_acc = hist6.history["val_accuracy"][28]
dl_prec = hist6.history["val_precision_m"][28]
dl_rec = hist6.history["val_recall_m"][28]
dl_f1 = hist6.history["val_f1_m"][28]

storeResults('TL - ResNet50',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [41]: x=hist6
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## Inception ResNetV2

```
In [42]: base_model = InceptionResNetV2(input_shape = (256,256,3), weights=None, include_top=True)

x1= Flatten()(base_model.output)
prediction1 = Dense(4, activation='softmax')(x1)
model7 = Model(inputs = base_model.inputs, outputs = prediction1)
model7.summary()
model7.compile(loss = 'categorical_crossentropy', optimizer='sgd', metrics=[ "accuracy",f1_m,precision_m, recall_m])
```

Model: "model\_6"

Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[(None, 256, 256, 3 0 )]	[ ]	
conv2d_94 (Conv2D)	(None, 127, 127, 32 864 )	[ 'input_7[0][0]' ]	
batch_normalization_94 (BatchN ormalization)	(None, 127, 127, 32 96 )	[ 'conv2d_94[0][0]' ]	
activation_94 (Activation)	(None, 127, 127, 32 0 )	[ 'batch_normalization_94[0][0]' ]	
conv2d_95 (Conv2D)	(None, 125, 125, 32 9216 )	[ 'activation_94[0][0]' ]	
... (10 more rows)	...	...	...

```
In [43]: hist7 = model7.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])
```

```
Epoch 1/50
2856/2856 [=====] - 341s 114ms/step - loss: 1.3843 - accuracy: 0.2726 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3816 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 323s 113ms/step - loss: 1.3841 - accuracy: 0.2775 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3804 - val_accuracy: 0.3593 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 323s 113ms/step - loss: 1.3829 - accuracy: 0.2803 - f1_m: 0.0000e+00 - precision_m: 0.0000e+00 - recall_m: 0.0000e+00 - val_loss: 1.3463 - val_accuracy: 0.3089 - val_f1_m: 0.0000e+00 - val_precision_m: 0.0000e+00 - val_recall_m: 0.0000e+00 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 323s 113ms/step - loss: 1.2896 - accuracy: 0.3846 - f1_m: 0.0770 - precision_m: 0.1155 - recall_m: 0.0578 - val_loss: 1.2871 - val_accuracy: 0.3585 - val_f1_m: 0.1118 - val_precision_m: 0.1646 - val_recall_m: 0.0854 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 324s 114ms/step - loss: 1.1685 - accuracy: 0.4393 - f1_m: 0.2465 - precision_m: 0.3697 - recall_m: 0.1849 - val_loss: 1.0830 - val_accuracy: 0.4859 - val_f1_m: 0.3039 - val_precision_m: 0.4146 - val_recall_m: 0.2485 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - 323s 113ms/step - loss: 1.1066 - accuracy: 0.4561 - f1_m: 0.2722 - precision_m: 0.4083 - recall_m: 0.2041 - val_loss: 1.1609 - val_accuracy: 0.4325 - val_f1_m: 0.2401 - val_precision_m: 0.3316 - val_recall_m: 0.1944 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 323s 113ms/step - loss: 1.0861 - accuracy: 0.4821 - f1_m: 0.3118 - precision_m: 0.4460 - recall_m: 0.2321 - val_loss: 1.1000 - val_accuracy: 0.4485 - val_f1_m: 0.2789 - val_precision_m: 0.3770 - val_recall_m: 0.2071 - lr: 0.0100
```



```
In [74]: dl_acc = hist7.history["val_accuracy"][22]
dl_prec = hist7.history["val_precision_m"][22]
dl_rec = hist7.history["val_recall_m"][22]
dl_f1 = hist7.history["val_f1_m"][22]
```

```
storeResults('TL - InceptionResNetV2',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [44]: x=hist7
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## Xception

```
In [45]: # Defining the pretrained base model
base = Xception(include_top=False, weights='imagenet', input_shape=(256,256,3))
x = base.output
x = GlobalAveragePooling2D()(x)
# Defining the head of the model where the prediction is conducted
head = Dense(4, activation='softmax')(x)
# Combining base and head
model8 = Model(inputs=base.input, outputs=head)
model8.compile(optimizer='sgd',
              loss = 'categorical_crossentropy',
              metrics=['accuracy',f1_m,precision_m, recall_m])
model8.summary()
```

Model: "model\_7"

Layer (type)	Output Shape	Param #	Connected to
input_8 (InputLayer)	[None, 256, 256, 3 0 ]	0	[]
block1_conv1 (Conv2D)	(None, 127, 127, 32 864 )	864	['input_8[0][0]']
block1_conv1_bn (BatchNormaliz ation)	(None, 127, 127, 32 128 )	128	['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, 127, 127, 32 0 )	0	['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, 125, 125, 64 18432 )	18432	['block1_conv1_act[0][0]']
... (11 more rows)	... (11 more rows)	... (11 more rows)	... (11 more rows)

```
In [46]: hist8 = model8.fit(train_set, epochs=50, validation_data=test_set,steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])
```

```
Epoch 1/50
2856/2856 [=====] - 86s 29ms/step - loss: 0.5615 - accuracy: 0.7959 - f1_m: 0.7437 - precision_m: 0.8046 - recall_m: 0.7132 - val_loss: 0.1866 - val_accuracy: 0.9474 - val_f1_m: 0.9418 - val_precision_m: 0.9520 - val_recall_m: 0.9367 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 81s 28ms/step - loss: 0.1618 - accuracy: 0.9464 - f1_m: 0.9442 - precision_m: 0.9561 - recall_m: 0.9382 - val_loss: 0.1152 - val_accuracy: 0.9649 - val_f1_m: 0.9652 - val_precision_m: 0.9688 - val_recall_m: 0.9634 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 81s 28ms/step - loss: 0.0840 - accuracy: 0.9730 - f1_m: 0.9716 - precision_m: 0.9767 - recall_m: 0.9690 - val_loss: 0.0529 - val_accuracy: 0.9855 - val_f1_m: 0.9860 - val_precision_m: 0.9870 - val_recall_m: 0.9855 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 80s 28ms/step - loss: 0.0467 - accuracy: 0.9856 - f1_m: 0.9854 - precision_m: 0.9862 - recall_m: 0.9849 - val_loss: 0.0993 - val_accuracy: 0.9687 - val_f1_m: 0.9688 - val_precision_m: 0.9688 - val_recall_m: 0.9688 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 80s 28ms/step - loss: 0.0334 - accuracy: 0.9907 - f1_m: 0.9904 - precision_m: 0.9911 - recall_m: 0.9900 - val_loss: 0.0606 - val_accuracy: 0.9825 - val_f1_m: 0.9822 - val_precision_m: 0.9832 - val_recall_m: 0.9817 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - 81s 28ms/step - loss: 0.0278 - accuracy: 0.9918 - f1_m: 0.9916 - precision_m: 0.9923 - recall_m: 0.9912 - val_loss: 0.0385 - val_accuracy: 0.9886 - val_f1_m: 0.9881 - val_precision_m: 0.9886 - val_recall_m: 0.9878 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 80s 28ms/step - loss: 0.0242 - accuracy: 0.9928 - f1_m: 0.9926 - precision_m: 0.9931 - recall_m: 0.9924 - val_loss: 0.0355 - val_accuracy: 0.9891 - val_f1_m: 0.9889 - val_precision_m: 0.9891 - val_recall_m: 0.9889 - lr: 0.0100
```

```
In [47]: dl_acc = hist8.history["val_accuracy"][-1]
dl_prec = hist8.history["val_precision_m"][-1]
dl_rec = hist8.history["val_recall_m"][-1]
dl_f1 = hist8.history["val_f1_m"][-1]


storeResults('TL - Xception',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [48]: model8.save('models/xception.h5')
```

```
In [48]: x=hist8
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## NASNetMobile

```
In [49]: # Defining the pretrained base model
base = NASNetMobile(include_top=False, weights='imagenet', input_shape=(256,256,3))
x = base.output
x = GlobalAveragePooling2D()(x)
# Defining the head of the model where the prediction is conducted
head = Dense(4, activation='softmax')(x)
# Combining base and head
model9 = Model(inputs=base.input, outputs=head)
model9.compile(optimizer='sgd',
              loss = 'categorical_crossentropy',
              metrics=["accuracy",f1_m,precision_m, recall_m])
model9.summary()

Model: "model_8"
-----
```

Layer (type)	Output Shape	Param #	Connected to
input_9 (InputLayer)	[None, 256, 256, 3 0 ]		
stem_conv1 (Conv2D)	(None, 127, 127, 32 864 )		['input_9[0][0]']
stem_bn1 (BatchNormalization)	(None, 127, 127, 32 128 )		['stem_conv1[0][0]']
activation_297 (Activation)	(None, 127, 127, 32 0 )		['stem_bn1[0][0]']
reduction_conv_1_stem_1 (Conv2D)	(None, 127, 127, 11 352 )		['activation_297[0][0]']
.....	.....	.....	.....

```
In [50]: hist9 = model9.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])

Epoch 1/50
2856/2856 [=====] - 289s 95ms/step - loss: 0.4118 - accuracy: 0.8503 - f1_m: 0.8386 - precision_m: 0.8708 - recall_m: 0.8225 - val_loss: 0.3290 - val_accuracy: 0.8902 - val_f1_m: 0.8786 - val_precision_m: 0.9040 - val_recall_m: 0.8659 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 264s 92ms/step - loss: 0.1619 - accuracy: 0.9454 - f1_m: 0.9445 - precision_m: 0.9519 - recall_m: 0.9408 - val_loss: 0.4295 - val_accuracy: 0.8459 - val_f1_m: 0.8280 - val_precision_m: 0.8636 - val_recall_m: 0.8102 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.1051 - accuracy: 0.9610 - f1_m: 0.9611 - precision_m: 0.9648 - recall_m: 0.9592 - val_loss: 0.1611 - val_accuracy: 0.9504 - val_f1_m: 0.9512 - val_precision_m: 0.9543 - val_recall_m: 0.9497 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0669 - accuracy: 0.9792 - f1_m: 0.9789 - precision_m: 0.9809 - recall_m: 0.9779 - val_loss: 0.1048 - val_accuracy: 0.9664 - val_f1_m: 0.9665 - val_precision_m: 0.9680 - val_recall_m: 0.9657 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0508 - accuracy: 0.9832 - f1_m: 0.9831 - precision_m: 0.9842 - recall_m: 0.9825 - val_loss: 0.0911 - val_accuracy: 0.9718 - val_f1_m: 0.9688 - val_precision_m: 0.9764 - val_recall_m: 0.9649 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - 264s 92ms/step - loss: 0.0426 - accuracy: 0.9860 - f1_m: 0.9859 - precision_m: 0.9867 - recall_m: 0.9855 - val_loss: 0.2065 - val_accuracy: 0.9359 - val_f1_m: 0.9357 - val_precision_m: 0.9413 - val_recall_m: 0.9329 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0346 - accuracy: 0.9898 - f1_m: 0.9898 - precision_m: 0.9902 - recall_m: 0.9897 - val_loss: 0.0690 - val_accuracy: 0.9786 - val_f1_m: 0.9776 - val_precision_m: 0.9802 - val_recall_m: 0.9764 - lr: 0.0100
Epoch 8/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0212 - accuracy: 0.9932 - f1_m: 0.9931 - precision_m: 0.9935 - recall_m: 0.9928 - val_loss: 0.0716 - val_accuracy: 0.9832 - val_f1_m: 0.9835 - val_precision_m: 0.9840 - val_recall_m: 0.9832 - lr: 0.0100
Epoch 9/50
2856/2856 [=====] - 264s 93ms/step - loss: 0.0294 - accuracy: 0.9905 - f1_m: 0.9904 - precision_m: 0.9907 - recall_m: 0.9902 - val_loss: 0.0536 - val_accuracy: 0.9840 - val_f1_m: 0.9842 - val_precision_m: 0.9848 - val_recall_m: 0.9840 - lr: 0.0100
Epoch 10/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0200 - accuracy: 0.9932 - f1_m: 0.9931 - precision_m: 0.9933 - recall_m: 0.9930 - val_loss: 0.1182 - val_accuracy: 0.9680 - val_f1_m: 0.9688 - val_precision_m: 0.9703 - val_recall_m: 0.9680 - lr: 0.0100
Epoch 11/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0247 - accuracy: 0.9912 - f1_m: 0.9915 - precision_m: 0.9921 - recall_m: 0.9912 - val_loss: 0.0503 - val_accuracy: 0.9840 - val_f1_m: 0.9848 - val_precision_m: 0.9863 - val_recall_m: 0.9840 - lr: 0.0100
Epoch 12/50
2856/2856 [=====] - ETA: 0s - loss: 0.0124 - accuracy: 0.9958 - f1_m: 0.9958 - precision_m: 0.9958 - recall_m: 0.9958
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.002999999329447745.
2856/2856 [=====] - 265s 93ms/step - loss: 0.0124 - accuracy: 0.9958 - f1_m: 0.9958 - precision_m: 0.9958 - recall_m: 0.9958 - val_loss: 0.1145 - val_accuracy: 0.9664 - val_f1_m: 0.9654 - val_precision_m: 0.9665 - val_recall_m: 0.9649 - lr: 0.0100
Epoch 13/50
2856/2856 [=====] - 264s 93ms/step - loss: 0.0073 - accuracy: 0.9977 - f1_m: 0.9977 - precision_m: 0.9979 - recall_m: 0.9975 - val_loss: 0.0301 - val_accuracy: 0.9939 - val_f1_m: 0.9942 - val_precision_m: 0.9947 - val_recall_m: 0.9939 - lr: 0.0030
Epoch 14/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0041 - accuracy: 0.9988 - f1_m: 0.9987 - precision_m: 0.9988 - recall_m: 0.9986 - val_loss: 0.0247 - val_accuracy: 0.9939 - val_f1_m: 0.9942 - val_precision_m: 0.9947 - val_recall_m: 0.9939 - lr: 0.0030
Epoch 15/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0024 - accuracy: 0.9995 - f1_m: 0.9995 - precision_m: 0.9995 - recall_m: 0.9995 - val_loss: 0.0223 - val_accuracy: 0.9931 - val_f1_m: 0.9931 - val_precision_m: 0.9931 - val_recall_m: 0.9931 - lr: 0.0030
Epoch 16/50
2856/2856 [=====] - ETA: 0s - loss: 0.0017 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996
Epoch 16: ReduceLROnPlateau reducing learning rate to 0.000900000078231095.
2856/2856 [=====] - 265s 93ms/step - loss: 0.0017 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996 - val_loss: 0.0297 - val_accuracy: 0.9916 - val_f1_m: 0.9916 - val_precision_m: 0.9916 - val_recall_m: 0.9916 - lr: 0.0030
Epoch 17/50
2856/2856 [=====] - 265s 93ms/step - loss: 9.9940e-04 - accuracy: 0.9998 - f1_m: 0.9998 - precision_m: 0.9998 - recall_m: 0.9998 - val_loss: 0.0245 - val_accuracy: 0.9939 - val_f1_m: 0.9939 - val_precision_m: 0.9939 - val_recall_m: 0.9939 - lr: 9.0000e-04
Epoch 18/50
2856/2856 [=====] - 264s 93ms/step - loss: 0.0022 - accuracy: 0.9995 - f1_m: 0.9995 - precision_m: 0.9995 - recall_m: 0.9995 - val_loss: 0.0218 - val_accuracy: 0.9931 - val_f1_m: 0.9931 - val_precision_m: 0.9931 - val_recall_m: 0.9931 - lr: 9.0000e-04
Epoch 19/50
2856/2856 [=====] - ETA: 0s - loss: 0.0019 - accuracy: 0.9993 - f1_m: 0.9993 - precision_m: 0.9993 - recall_m: 0.9993
Epoch 19: ReduceLROnPlateau reducing learning rate to 0.002699999536201356.
2856/2856 [=====] - 265s 93ms/step - loss: 0.0019 - accuracy: 0.9993 - f1_m: 0.9993 - precision_m: 0.9993 - recall_m: 0.9993 - val_loss: 0.0220 - val_accuracy: 0.9939 - val_f1_m: 0.9934 - val_precision_m: 0.9939 - val_recall_m: 0.9931 - lr: 9.0000e-04
Epoch 20/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0018 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996 - val_loss: 0.0223 - val_accuracy: 0.9939 - val_f1_m: 0.9934 - val_precision_m: 0.9939 - val_recall_m: 0.9931 - lr: 2.7000e-04
Epoch 21/50
2856/2856 [=====] - 265s 93ms/step - loss: 0.0018 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996 - val_loss: 0.0222 - val_accuracy: 0.9947 - val_f1_m: 0.9942 - val_precision_m: 0.9947 - val_recall_m: 0.9939 - lr: 2.7000e-04
Epoch 22/50
2856/2856 [=====] - 264s 93ms/step - loss: 0.0012 - accuracy: 0.9998 - f1_m: 0.9998 - precision_m: 0.9998 - recall_m: 0.9998 - val_loss: 0.0218 - val_accuracy: 0.9947 - val_f1_m: 0.9942 - val_precision_m: 0.9947 - val_recall_m: 0.9939 - lr: 2.7000e-04
Epoch 23/50
2856/2856 [=====] - ETA: 0s - loss: 0.0014 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996Restoring model weights from the end of the best epoch: 13.
2856/2856 [=====] - 265s 93ms/step - loss: 0.0014 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9996 - recall_m: 0.9996 - val_loss: 0.0220 - val_accuracy: 0.9947 - val_f1_m: 0.9942 - val_precision_m: 0.9947 - val_recall_m: 0.9939 - lr: 2.7000e-04
Epoch 23: early stopping
```

```
In [76]: dl_acc = hist9.history["val_accuracy"][22]
dl_prec = hist9.history["val_precision_m"][22]
dl_rec = hist9.history["val_recall_m"][22]
dl_f1 = hist9.history["val_f1_m"][22]



storeResults('TL - NASNetMobile',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [51]: model9.save('models/nasnetmobile.h5')
```

```
In [52]: x=hist9
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=16)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## NASNetLarge

```
In [53]: # Defining the pretrained base model
base = NASNetLarge(include_top=False, weights='imagenet', input_shape=(256,256,3))
x = base.output
x = GlobalAveragePooling2D()(x)
# Defining the head of the model where the prediction is conducted
head = Dense(4, activation='softmax')(x)
# Combining base and head
model10 = Model(inputs=base.input, outputs=head)
model10.compile(optimizer='sgd',
                 loss = 'categorical_crossentropy',
                 metrics=['accuracy',f1_m,precision_m, recall_m])
model10.summary()
```

Model: "model\_9"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_10 (InputLayer)	[None, 256, 256, 3 0 ]	0	[]
stem_conv1 (Conv2D)	(None, 127, 127, 96 2592 )	2592	['input_10[0][0]']
stem_bn1 (BatchNormalization)	(None, 127, 127, 96 384 )	384	['stem_conv1[0][0]']
activation_485 (Activation)	(None, 127, 127, 96 0 )	0	['stem_bn1[0][0]']
reduction_conv_1_stem_1 (Conv2D)	(None, 127, 127, 42 4032 )	4032	['activation_485[0][0]']
... (many more layers)	...	...	...

```
In [54]: hist10 = model10.fit(train_set, epochs=50, validation_data=test_set, steps_per_epoch=len(train_set), validation_steps=len(test_set), callbacks=[learning_rate_reduction, early_stop])
```

Epoch 1/50  
2856/2856 [=====] - 395s 130ms/step - loss: 0.3068 - accuracy: 0.8922 - f1\_m: 0.8748 - precision\_m: 0.9018 - recall\_m: 0.8613 - val\_loss: 0.1955 - val\_accuracy: 0.9390 - val\_f1\_m: 0.9400 - val\_precision\_m: 0.9436 - val\_recall\_m: 0.9383 - lr: 0.0100  
Epoch 2/50  
2856/2856 [=====] - 366s 128ms/step - loss: 0.0828 - accuracy: 0.9723 - f1\_m: 0.9723 - precision\_m: 0.9764 - recall\_m: 0.9702 - val\_loss: 0.0982 - val\_accuracy: 0.9687 - val\_f1\_m: 0.9690 - val\_precision\_m: 0.9756 - val\_recall\_m: 0.9657 - lr: 0.0100  
Epoch 3/50  
2856/2856 [=====] - 366s 128ms/step - loss: 0.0473 - accuracy: 0.9842 - f1\_m: 0.9833 - precision\_m: 0.9853 - recall\_m: 0.9823 - val\_loss: 0.0660 - val\_accuracy: 0.9817 - val\_f1\_m: 0.9817 - val\_precision\_m: 0.9863 - val\_recall\_m: 0.9794 - lr: 0.0100  
Epoch 4/50  
2856/2856 [=====] - 366s 128ms/step - loss: 0.0376 - accuracy: 0.9886 - f1\_m: 0.9889 - precision\_m: 0.9897 - recall\_m: 0.9884 - val\_loss: 0.0562 - val\_accuracy: 0.9870 - val\_f1\_m: 0.9876 - val\_precision\_m: 0.9901 - val\_recall\_m: 0.9863 - lr: 0.0100  
Epoch 5/50  
2856/2856 [=====] - 366s 128ms/step - loss: 0.0223 - accuracy: 0.9935 - f1\_m: 0.9935 - precision\_m: 0.9939 - recall\_m: 0.9933 - val\_loss: 0.0388 - val\_accuracy: 0.9916 - val\_f1\_m: 0.9916 - val\_precision\_m: 0.9916 - val\_recall\_m: 0.9916 - lr: 0.0100  
Epoch 6/50  
2856/2856 [=====] - 365s 128ms/step - loss: 0.0140 - accuracy: 0.9961 - f1\_m: 0.9960 - precision\_m: 0.9961 - recall\_m: 0.9960 - val\_loss: 0.0341 - val\_accuracy: 0.9916 - val\_f1\_m: 0.9916 - val\_precision\_m: 0.9916 - val\_recall\_m: 0.9916 - lr: 0.0100  
Epoch 7/50  
2856/2856 [=====] - 365s 128ms/step - loss: 0.0140 - accuracy: 0.9961 - f1\_m: 0.9960 - precision\_m: 0.9961 - recall\_m: 0.9960 - val\_loss: 0.0341 - val\_accuracy: 0.9916 - val\_f1\_m: 0.9916 - val\_precision\_m: 0.9916 - val\_recall\_m: 0.9916 - lr: 0.0100

```
In [77]: dl_acc = hist10.history["val_accuracy"][27]
dl_prec = hist10.history["val_precision_m"][27]
dl_rec = hist10.history["val_recall_m"][27]
dl_f1 = hist10.history["val_f1_m"][27]



storeResults('TL - NASNetLarge',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [55]: model10.save('models/nasnetlarge.h5')
```

```
In [56]: x=hist10
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## IVX16

```
In [57]: from tensorflow.keras.models import load_model
from tensorflow.keras.layers import Average
import efficientnet.keras
```

```
In [58]: def ensemble():

    model_1 = load_model('models/inceptionv3.h5', compile=False)
    model_1 = Model(inputs = model_1.inputs, outputs = model_1.outputs, name = 'InceptionV3')

    model_2 = load_model('models/xception.h5', compile=False)
    model_2 = Model(inputs = model_2.inputs, outputs = model_2.outputs, name = 'Xception')

    model_3 = load_model('models/vgg16.h5', compile=False)
    model_3 = Model(inputs = model_3.inputs, outputs = model_3.outputs, name = 'VGG16')

    models = [model_1, model_2, model_3]

    models_input = Input(shape =(256,256,3))
    models_output = [model(models_input) for model in models]

    ensemble_output = Average()(models_output)

    simple_average = Model(inputs = models_input, outputs = ensemble_output, name = 'IVX16')

    return simple_average
```

```
In [59]: model = ensemble()
model.compile(optimizer='sgd',
              loss = 'categorical_crossentropy',
              metrics=["accuracy",f1_m,precision_m, recall_m])
model.summary()
```

Model: "IVX16"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_11 (InputLayer)	[(None, 256, 256, 3 0 )]	2	[]
InceptionV3 (Functional)	(None, 4)	23855788	['input_11[0][0]']
Xception (Functional)	(None, 4)	20869676	['input_11[0][0]']
average (Average)	(None, 4)	0	['InceptionV3[0][0]', 'Xception[0][0]']
<hr/>			
Total params:	44,725,464		
Trainable params:	44,636,504		
Non-trainable params:	88,960		

```
In [60]: history = model.fit(
    train_set,
    epochs=50,
    validation_data=test_set,callbacks=[learning_rate_reduction, early_stop])
```

```
Epoch 1/50
2856/2856 [=====] - 224s 75ms/step - loss: 0.3555 - accuracy: 0.9996 - f1_m: 0.9992 - precision_m: 1.0000 - recall_m: 0.9988 - val_loss: 0.3732 - val_accuracy: 0.9954 - val_f1_m: 0.9936 - val_precision_m: 0.9962 - val_recall_m: 0.9924 - lr: 0.0100
Epoch 2/50
2856/2856 [=====] - 214s 75ms/step - loss: 0.3526 - accuracy: 1.0000 - f1_m: 1.0000 - precision_m: 1.0000 - recall_m: 1.0000 - val_loss: 0.3717 - val_accuracy: 0.9954 - val_f1_m: 0.9949 - val_precision_m: 0.9970 - val_recall_m: 0.9939 - lr: 0.0100
Epoch 3/50
2856/2856 [=====] - 213s 75ms/step - loss: 0.3463 - accuracy: 0.9996 - f1_m: 0.9996 - precision_m: 0.9998 - recall_m: 0.9995 - val_loss: 0.3512 - val_accuracy: 0.9969 - val_f1_m: 0.9931 - val_precision_m: 0.9977 - val_recall_m: 0.9909 - lr: 0.0100
Epoch 4/50
2856/2856 [=====] - 213s 75ms/step - loss: 0.3448 - accuracy: 1.0000 - f1_m: 0.9999 - precision_m: 1.0000 - recall_m: 0.9998 - val_loss: 0.3509 - val_accuracy: 0.9954 - val_f1_m: 0.9936 - val_precision_m: 0.9962 - val_recall_m: 0.9924 - lr: 0.0100
Epoch 5/50
2856/2856 [=====] - 214s 75ms/step - loss: 0.3466 - accuracy: 0.9998 - f1_m: 0.9998 - precision_m: 1.0000 - recall_m: 0.9996 - val_loss: 0.3338 - val_accuracy: 0.9962 - val_f1_m: 0.9954 - val_precision_m: 0.9970 - val_recall_m: 0.9947 - lr: 0.0100
Epoch 6/50
2856/2856 [=====] - ETA: 0s - loss: 0.3419 - accuracy: 0.9998 - f1_m: 0.9994 - precision_m: 1.0000 - recall_m: 0.9991
Epoch 6: ReduceLROnPlateau reducing learning rate to 0.002999999329447745.
2856/2856 [=====] - 213s 74ms/step - loss: 0.3419 - accuracy: 0.9998 - f1_m: 0.9994 - precision_m: 1.0000 - recall_m: 0.9991 - val_loss: 0.3273 - val_accuracy: 0.9931 - val_f1_m: 0.9942 - val_precision_m: 0.9962 - val_recall_m: 0.9931 - lr: 0.0100
Epoch 7/50
2856/2856 [=====] - 213s 75ms/step - loss: 0.3410 - accuracy: 0.9998 - f1_m: 0.9999 - precision_m: 1.0000 - recall_m: 0.9998 - val_loss: 0.3528 - val_accuracy: 0.9931 - val_f1_m: 0.9934 - val_precision_m: 0.9970 - val_recall_m: 0.9916 - lr: 0.0030
Epoch 8/50
2856/2856 [=====] - 214s 75ms/step - loss: 0.3399 - accuracy: 0.9998 - f1_m: 0.9997 - precision_m: 0.9998 - recall_m: 0.9996 - val_loss: 0.3472 - val_accuracy: 0.9947 - val_f1_m: 0.9944 - val_precision_m: 0.9954 - val_recall_m: 0.9939 - lr: 0.0030
Epoch 9/50
2856/2856 [=====] - ETA: 0s - loss: 0.3391 - accuracy: 1.0000 - f1_m: 0.9996 - precision_m: 1.0000 - recall_m: 0.9995
Epoch 9: ReduceLROnPlateau reducing learning rate to 0.000900000078231095.
2856/2856 [=====] - 213s 75ms/step - loss: 0.3391 - accuracy: 1.0000 - f1_m: 0.9996 - precision_m: 1.0000 - recall_m: 0.9995 - val_loss: 0.3492 - val_accuracy: 0.9947 - val_f1_m: 0.9944 - val_precision_m: 0.9970 - val_recall_m: 0.9931 - lr: 0.0030
Epoch 10/50
2856/2856 [=====] - 214s 75ms/step - loss: 0.3361 - accuracy: 1.0000 - f1_m: 0.9996 - precision_m: 1.0000 - recall_m: 0.9995 - val_loss: 0.3515 - val_accuracy: 0.9924 - val_f1_m: 0.9931 - val_precision_m: 0.9962 - val_recall_m: 0.9916 - lr: 9.0000e-04
Epoch 11/50
2856/2856 [=====] - 214s 75ms/step - loss: 0.3361 - accuracy: 1.0000 - f1_m: 0.9996 - precision_m: 1.0000 - recall_m: 0.9995 - val_loss: 0.3481 - val_accuracy: 0.9954 - val_f1_m: 0.9944 - val_precision_m: 0.9970 - val_recall_m: 0.9931 - lr: 9.0000e-04
Epoch 12/50
2856/2856 [=====] - ETA: 0s - loss: 0.3352 - accuracy: 1.0000 - f1_m: 1.0000 - precision_m: 1.0000 - recall_m: 1.0000
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.0002699999536201356.
2856/2856 [=====] - 213s 75ms/step - loss: 0.3352 - accuracy: 1.0000 - f1_m: 1.0000 - precision_m: 1.0000 - recall_m: 1.0000 - val_loss: 0.3376 - val_accuracy: 0.9947 - val_f1_m: 0.9954 - val_precision_m: 0.9970 - val_recall_m: 0.9947 - lr: 9.0000e-04
Epoch 13/50
2856/2856 [=====] - ETA: 0s - loss: 0.3369 - accuracy: 1.0000 - f1_m: 0.9998 - precision_m: 1.0000 - recall_m: 0.9996Restoring model weights from the end of the best epoch: 3.
2856/2856 [=====] - 214s 75ms/step - loss: 0.3369 - accuracy: 1.0000 - f1_m: 0.9998 - precision_m: 1.0000 - recall_m: 0.9996 - val_loss: 0.3410 - val_accuracy: 0.9947 - val_f1_m: 0.9936 - val_precision_m: 0.9962 - val_recall_m: 0.9924 - lr: 2.7000e-04
Epoch 13: early stopping
```



```
In [78]: dl_acc = history.history["val_accuracy"][12]
dl_prec = history.history["val_precision_m"][12]
dl_rec = history.history["val_recall_m"][12]
dl_f1 = history.history["val_f1_m"][12]

storeResults('TL - IVX16',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [61]: x=history
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=16)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## Extension

```
In [62]: def ensemble():

    model_1 = load_model('models/nasnetlarge.h5', compile=False)
    model_1 = Model(inputs = model_1.inputs, outputs = model_1.outputs, name = 'NASNetLarge')

    model_2 = load_model('models/nasnetmobile.h5', compile=False)
    model_2 = Model(inputs = model_2.inputs, outputs = model_2.outputs, name = 'NASNetMobile')

    models = [model_1, model_2]

    models_input = Input(shape =(256,256,3))
    models_output = [model(models_input) for model in models]

    ensemble_output = Average()(models_output)

    simple_average = Model(inputs = models_input, outputs = ensemble_output, name = 'Extension')

    return simple_average
```

```
In [63]: ext = ensemble()
ext.compile(optimizer='sgd',
            loss = 'categorical_crossentropy',
            metrics=["accuracy",f1_m,precision_m, recall_m])
ext.summary()
```

Model: "Extension"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_12 (InputLayer)	[(None, 256, 256, 3 0 )]	0	[]
NASNetLarge (Functional)	(None, 4)	84932950	['input_12[0][0]']
NASNetMobile (Functional)	(None, 4)	4273944	['input_12[0][0]']
average_1 (Average)	(None, 4)	0	['NASNetLarge[0][0]', 'NASNetMobile[0][0]']
<hr/>			

```
Total params: 89,206,894
Trainable params: 88,973,488
Non-trainable params: 233,406
```

---

```
In [64]: history1 = ext.fit(
    train_set,
    epochs=50,
    validation_data=test_set, callbacks=[learning_rate_reduction, early_stop])
```

Epoch 1/50  
2856/2856 [=====] - 672s 220ms/step - loss: 0.0014 - accuracy: 0.9998 - f1\_m: 0.9998 - precision\_m: 0.9998 - recall\_m: 0.9998 - val\_loss: 0.0225 - val\_accuracy: 0.9947 - val\_f1\_m: 0.9952 - val\_precision\_m: 0.9962 - val\_recall\_m: 0.9947 - lr: 0.0100  
Epoch 2/50  
2856/2856 [=====] - 618s 216ms/step - loss: 9.3611e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0206 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 0.0100  
Epoch 3/50  
2856/2856 [=====] - 617s 216ms/step - loss: 6.8648e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0218 - val\_accuracy: 0.9962 - val\_f1\_m: 0.9962 - val\_precision\_m: 0.9962 - val\_recall\_m: 0.9962 - lr: 0.0100  
Epoch 4/50  
2856/2856 [=====] - 617s 216ms/step - loss: 7.8715e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0163 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9970 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9970 - lr: 0.0100  
Epoch 5/50  
2856/2856 [=====] - ETA: 0s - loss: 0.0012 - accuracy: 0.9996 - f1\_m: 0.9996 - precision\_m: 0.9996 - recall\_m: 0.9996  
Epoch 5: ReduceLROnPlateau reducing learning rate to 0.002999999329447745.  
2856/2856 [=====] - 617s 216ms/step - loss: 0.0012 - accuracy: 0.9996 - f1\_m: 0.9996 - precision\_m: 0.9996 - recall\_m: 0.9996 - val\_loss: 0.0235 - val\_accuracy: 0.9954 - val\_f1\_m: 0.9952 - val\_precision\_m: 0.9962 - val\_recall\_m: 0.9947 - lr: 0.0100  
Epoch 6/50  
2856/2856 [=====] - 617s 216ms/step - loss: 9.2013e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0203 - val\_accuracy: 0.9962 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 0.0030  
Epoch 7/50  
2856/2856 [=====] - 616s 216ms/step - loss: 6.1913e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0202 - val\_accuracy: 0.9962 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 0.0030  
Epoch 8/50  
2856/2856 [=====] - ETA: 0s - loss: 4.1802e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000  
Epoch 8: ReduceLROnPlateau reducing learning rate to 0.000900000078231095.  
2856/2856 [=====] - 616s 216ms/step - loss: 4.1802e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0191 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 0.0030  
Epoch 9/50  
2856/2856 [=====] - 616s 216ms/step - loss: 4.6883e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0186 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 9.0000e-04  
Epoch 10/50  
2856/2856 [=====] - 615s 215ms/step - loss: 5.7156e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0193 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 9.0000e-04  
Epoch 11/50  
2856/2856 [=====] - ETA: 0s - loss: 4.3507e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000  
Epoch 11: ReduceLROnPlateau reducing learning rate to 0.0002699999536201356.  
2856/2856 [=====] - 617s 216ms/step - loss: 4.3507e-04 - accuracy: 1.0000 - f1\_m: 1.0000 - precision\_m: 1.0000 - recall\_m: 1.0000 - val\_loss: 0.0192 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 9.0000e-04  
Epoch 12/50  
2856/2856 [=====] - ETA: 0s - loss: 6.1644e-04 - accuracy: 0.9998 - f1\_m: 0.9998 - precision\_m: 0.9998 - recall\_m: 0.9998Restoring model weights from the end of the best epoch: 2.  
2856/2856 [=====] - 617s 216ms/step - loss: 6.1644e-04 - accuracy: 0.9998 - f1\_m: 0.9998 - precision\_m: 0.9998 - recall\_m: 0.9998 - val\_loss: 0.0192 - val\_accuracy: 0.9969 - val\_f1\_m: 0.9959 - val\_precision\_m: 0.9970 - val\_recall\_m: 0.9954 - lr: 2.7000e-04  
Epoch 12: early stopping



```
In [81]: dl_acc = history1.history["val_accuracy"][11]
dl_prec = history1.history["val_precision_m"][11]
dl_rec = history1.history["val_recall_m"][11]
dl_f1 = history1.history["val_f1_m"][11]

storeResults('TL - Extension',dl_acc,dl_prec,dl_rec,dl_f1)
```

```
In [80]: x=history1
plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(x.history['loss'], label='Training Loss')
plt.plot(x.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(x.history['accuracy'], label='Training Accuracy')
plt.plot(x.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Optimizer : adam



## Comparison

```
In [82]: #creating dataframe
import pandas as pd
result = pd.DataFrame({ 'ML Model' : ML_Model,
                        'Accuracy' : accuracy,
                        'Precision': precision,
                        'Recall' : recall,
                        'F1-Score' : f1score,
})
})
```

```
In [83]: result
```

```
Out[83]:
```

	ML Model	Accuracy	Precision	Recall	F1-Score
0	VisionTransformer - SWIN	0.811	0.818	0.797	0.804
1	VisionTransformer - CCT	0.590	0.617	0.495	0.536
2	VisionTransformer - EANet	0.511	0.511	0.504	0.506
3	TL - VGG16	0.309	0.095	0.309	0.146
4	TL - VGG19	0.309	0.095	0.309	0.146
5	TL - InceptionV3	0.477	0.434	0.245	0.308
6	TL - ResNet50	0.665	0.689	0.463	0.539
7	TL - InceptionResNetV2	0.484	0.429	0.247	0.308
8	TL - Xception	0.995	0.995	0.995	0.995
9	TL - NASNetMobile	0.995	0.995	0.994	0.994
10	TL - NASNetLarge	0.994	0.994	0.994	0.994
11	TL - IVX16	0.995	0.996	0.992	0.994
12	TL - Extension	0.997	0.997	0.995	0.996

## Modelling

```
In [84]: ext.save('models/extension.h5')
```


## Graph

```
In [85]: classifier = ML_Model
y_pos = np.arange(len(classifier))
```

## Accuracy

```
In [86]: import matplotlib.pyplot as plt2
plt2.barch(y_pos, accuracy, align='center', alpha=0.5,color='blue')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Accuracy Score')
plt2.title('Classification Performance')
plt2.show()
```


Classification Performance



## Precision


```
In [87]: plt2.barch(y_pos, precision, align='center', alpha=0.5,color='red')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Precision Score')
plt2.title('Classification Performance')
plt2.show()
```

Classification Performance




## Recall

```
In [88]: plt2.barh(y_pos, recall, align='center', alpha=0.5,color='green')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Recall Score')
plt2.title('Classification Performance')
plt2.show()
```



## F1 Score

```
In [89]: plt2.barh(y_pos, f1score, align='center', alpha=0.5,color='navy')
plt2.yticks(y_pos, classifier)
plt2.xlabel('F1 Score')
plt2.title('Classification Performance')
plt2.show()
```



In [ ]: